

## ASSIGMENT 2

```
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import pandas as pd
import tensorflow as tf
import pickle
from tensorflow.keras.datasets import cifar10

# Load CIFAR-10 dataset
(train_images, train_labels), (test_images, test_labels) = cifar10.load_data()

# CIFAR-10 class names
class_names = ['Airplane', 'Automobile', 'Bird', 'Cat', 'Deer', 'Dog', 'Frog', 'Horse', 'Ship', 'Truck']

# Check if precomputed data is available
try:
    with open('cifar10_stats.pkl', 'rb') as f:
        basic_stats, channel_stats, mean_images = pickle.load(f)
        print("Loaded precomputed statistics and mean images from file.")
except FileNotFoundError:
    print("No precomputed data found. Calculating statistics...")

    # Calculate basic statistics about the images
    mean_pixel_value = train_images.mean()
    std_pixel_value = train_images.std()
    min_pixel_value = train_images.min()
    max_pixel_value = train_images.max()

    # Create a DataFrame for basic statistics
    basic_stats = pd.DataFrame({
        'Statistic': ['Mean', 'Standard Deviation', 'Min', 'Max'],
        'Pixel Value': [mean_pixel_value, std_pixel_value, min_pixel_value,
max_pixel_value]
    })

    # Calculate per-channel statistics
    mean_per_channel = train_images.mean(axis=(0, 1, 2))
    std_per_channel = train_images.std(axis=(0, 1, 2))

    # Create a DataFrame for per-channel statistics
    channel_stats = pd.DataFrame({
        'Channel': ['Red', 'Green', 'Blue'],
        'Mean Pixel Value': mean_per_channel,
        'Standard Deviation': std_per_channel
    })

    # Calculate mean images for each class
```

## ASSIGNMENT 2

```
mean_images = []
for i in range(10):
    mean_image = train_images[train_labels.ravel() == i].mean(axis=0)
    mean_images.append(mean_image)

# Save calculated statistics and mean images
with open('cifar10_stats.pkl', 'wb') as f:
    pickle.dump((basic_stats, channel_stats, mean_images), f)
print("Statistics and mean images saved to file.")

# Display the shape of the dataset
print(f'Train images shape: {train_images.shape}')
print(f'Train labels shape: {train_labels.shape}')
print(f'Test images shape: {test_images.shape}')
print(f'Test labels shape: {test_labels.shape}')

# Display basic statistics
print("Basic Statistics of Pixel Values in CIFAR-10 Training Set:")
print(basic_stats)

# Display per-channel statistics
print("\nPer-Channel Statistics of Pixel Values in CIFAR-10 Training Set:")
print(channel_stats)

# Function to display a grid of images
def plot_image_grid(images, labels, class_names, rows=4, cols=4):
    fig, axes = plt.subplots(rows, cols, figsize=(12, 12))
    axes = axes.ravel()
    for i in np.arange(0, rows * cols):
        axes[i].imshow(images[i])
        axes[i].set_title(class_names[int(labels[i])])
        axes[i].axis('off')
    plt.subplots_adjust(hspace=0.5)
    plt.show()

# Plot a 4x4 grid of images from the training dataset
plot_image_grid(train_images, train_labels, class_names)

# Show an individual image with its label
def plot_single_image(index):
    plt.figure(figsize=(4, 4))
    plt.imshow(train_images[index])
    plt.title(f'Label: {class_names[int(train_labels[index])]}')
    plt.axis('off')
    plt.show()

# Plot a single image (e.g., image at index 0)
plot_single_image(0)
```

## ASSIGNMENT 2

```
# Display the distribution of classes in the training dataset
plt.figure(figsize=(12, 6))
sns.countplot(x=train_labels.ravel(), palette='viridis')
plt.xlabel('Class')
plt.ylabel('Frequency')
plt.xticks(np.arange(len(class_names)), class_names, rotation=45)
plt.title('Class Distribution in CIFAR-10 Training Set')
plt.show()

# Plot the distribution of pixel values
plt.figure(figsize=(10, 5))
sns.histplot(train_images.ravel(), bins=50, color='blue', kde=True)
plt.xlabel('Pixel Value')
plt.ylabel('Frequency')
plt.title('Distribution of Pixel Values in Training Set')
plt.show()

# Plot the mean image of each class
fig, axes = plt.subplots(2, 5, figsize=(15, 6))
axes = axes.ravel()
for i in range(10):
    axes[i].imshow(mean_images[i].astype(np.uint8))
    axes[i].set_title(class_names[i])
    axes[i].axis('off')
plt.subplots_adjust(hspace=0.5)
plt.show()

# Display class counts as a bar plot
plt.figure(figsize=(10, 6))
sns.barplot(x=class_names, y=class_counts, palette='viridis')
plt.xlabel('Class')
plt.ylabel('Count')
plt.title('Class Frequency in CIFAR-10 Training Set')
plt.xticks(rotation=45)
plt.show()
```

## ASSIGMENT 2

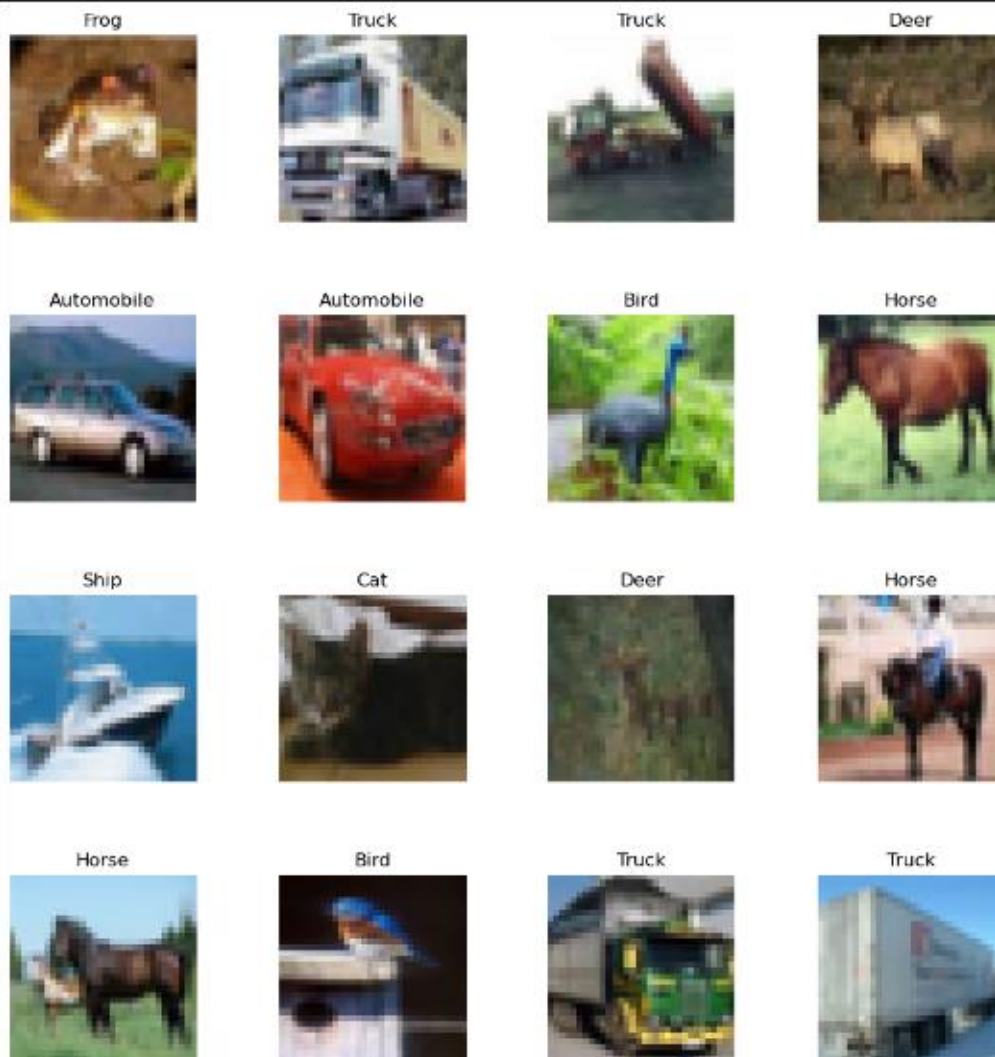
```
Loaded precomputed statistics and mean images from file.
Train images shape: (50000, 32, 32, 3)
Train labels shape: (50000, 1)
Test images shape: (10000, 32, 32, 3)
Test labels shape: (10000, 1)
```

Basic Statistics of Pixel Values in CIFAR-10 Training Set:




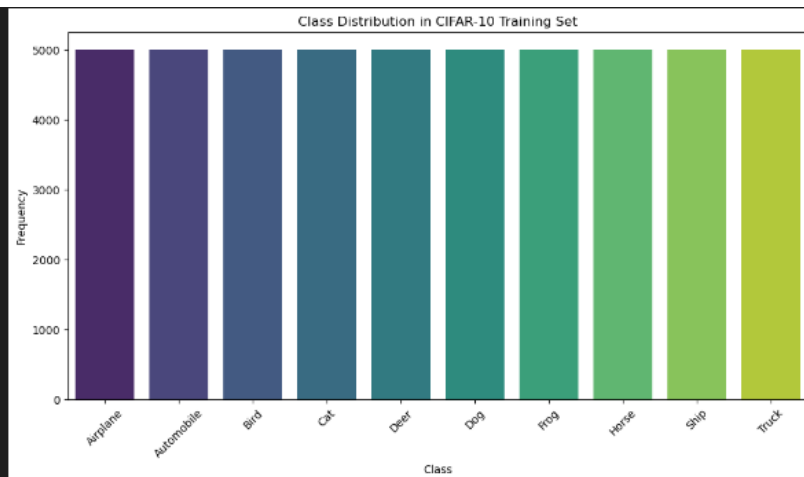
|   | Statistic          | Pixel Value |
|---|--------------------|-------------|
| 0 | Mean               | 120.707565  |
| 1 | Standard Deviation | 64.150876   |
| 2 | Min                | 0.000000    |
| 3 | Max                | 255.000000  |

Per-Channel Statistics of Pixel Values in CIFAR-10 Training Set:

| Channel | Mean Pixel Value | Standard Deviation |
|---------|------------------|--------------------|
| 0 Red   | 125.306918       | 62.993219          |
| 1 Green | 122.950394       | 62.088708          |
| 2 Blue  | 113.865383       | 66.704900          |



## ASSIGNMENT 2



## ASSIGMENT 2

```
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import pickle
from tensorflow.keras.datasets import cifar10
from sklearn.decomposition import PCA

# Load CIFAR-10 dataset
(train_images, train_labels), (test_images, test_labels) = cifar10.load_data()

# CIFAR-10 class names
class_names = ['Airplane', 'Automobile', 'Bird', 'Cat', 'Deer', 'Dog', 'Frog', 'Horse', 'Ship', 'Truck']

# Check if precomputed data is available
try:
    with open('cifar10_stats.pkl', 'rb') as f:
        basic_stats, channel_stats, mean_images = pickle.load(f)
        print("Loaded precomputed statistics and mean images from file.")
except FileNotFoundError:
    print("No precomputed data found. Calculating statistics...")

    # Calculate basic statistics about the images
    mean_pixel_value = train_images.mean()
    std_pixel_value = train_images.std()
    min_pixel_value = train_images.min()
    max_pixel_value = train_images.max()

    # Create a DataFrame for basic statistics
    basic_stats = pd.DataFrame({
        'Statistic': ['Mean', 'Standard Deviation', 'Min', 'Max'],
        'Pixel Value': [mean_pixel_value, std_pixel_value, min_pixel_value,
max_pixel_value]
    })

    # Calculate per-channel statistics
    mean_per_channel = train_images.mean(axis=(0, 1, 2))
    std_per_channel = train_images.std(axis=(0, 1, 2))

    # Create a DataFrame for per-channel statistics
    channel_stats = pd.DataFrame({
        'Channel': ['Red', 'Green', 'Blue'],
        'Mean Pixel Value': mean_per_channel,
        'Standard Deviation': std_per_channel
    })

    # Calculate mean images for each class
```

## ASSIGNMENT 2

```
mean_images = []
for i in range(10):
    mean_image = train_images[train_labels.ravel() == i].mean(axis=0)
    mean_images.append(mean_image)

# Save calculated statistics and mean images
with open('cifar10_stats.pkl', 'wb') as f:
    pickle.dump((basic_stats, channel_stats, mean_images), f)
print("Statistics and mean images saved to file.")

# Display the shape of the dataset
print(f'Train images shape: {train_images.shape}')
print(f'Train labels shape: {train_labels.shape}')
print(f'Test images shape: {test_images.shape}')
print(f'Test labels shape: {test_labels.shape}')

# Display basic statistics
print("Basic Statistics of Pixel Values in CIFAR-10 Training Set:")
print(basic_stats)

# Display per-channel statistics
print("\nPer-Channel Statistics of Pixel Values in CIFAR-10 Training Set:")
print(channel_stats)

# Perform PCA to reduce dimensionality of images for visualization purposes
n_samples = 1000 # Use a subset of the data for faster computation
flat_images = train_images[:n_samples].reshape(n_samples, -1)

pca = PCA(n_components=2)
principal_components = pca.fit_transform(flat_images)

# Create a DataFrame with PCA components and labels
pca_df = pd.DataFrame({
    'PCA1': principal_components[:, 0],
    'PCA2': principal_components[:, 1],
    'Label': train_labels[:n_samples].ravel()
})

# Plot PCA components to visualize class separability
plt.figure(figsize=(10, 8))
sns.scatterplot(x='PCA1', y='PCA2', hue='Label', palette='tab10', data=pca_df,
legend='full', alpha=0.7)
plt.title('PCA of CIFAR-10 Training Set')
plt.xlabel('PCA Component 1')
plt.ylabel('PCA Component 2')
plt.legend(class_names)
plt.show()
```

## ASSIGMENT 2

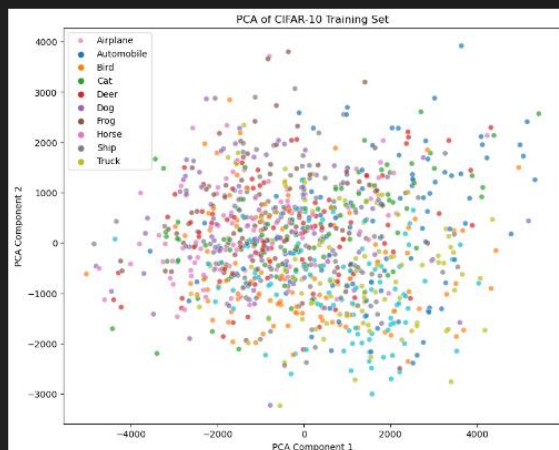
```
# Filter data by specific classes for further analysis
# Example: Filtering only 'Airplane' and 'Automobile' classes
filtered_indices = np.where((train_labels == 0) | (train_labels == 1))[0]
filtered_images = train_images[filtered_indices]
filtered_labels = train_labels[filtered_indices]

# Display filtered dataset information
print(f'Filtered images shape: {filtered_images.shape}')
print(f'Filtered labels shape: {filtered_labels.shape}')


# Display the distribution of the filtered classes
plt.figure(figsize=(8, 4))
sns.countplot(x=filtered_labels.ravel(), palette='viridis')
plt.xlabel('Class')
plt.ylabel('Frequency')
plt.xticks([0, 1], ['Airplane', 'Automobile'], rotation=45)
plt.title('Class Distribution in Filtered CIFAR-10 Dataset')
plt.show()
```

```
Loaded precomputed statistics and mean images from file.
Train images shape: (50000, 32, 32, 3)
Train labels shape: (50000, 1)
Test images shape: (10000, 32, 32, 3)
Test labels shape: (10000, 1)
Basic Statistics of Pixel Values in CIFAR-10 Training Set:
Statistic Pixel Value
0 Mean 125.787665
1 Standard Deviation 64.150676
2 Min 0.000000
3 Max 255.000000
```

```
Per-Channel Statistics of Pixel Values in CIFAR-10 Training Set:
Channel Mean Pixel Value Standard Deviation
0 Red 125.306918 62.993219
1 Green 122.958394 62.085708
2 Blue 113.865383 66.704500
```



```
Filtered images shape: (10000, 32, 32, 3)
Filtered labels shape: (10000, 1)
```





## ASSIGNMENT 2

```
import matplotlib.pyplot as plt
import numpy as np
import seaborn as sns
from sklearn.metrics import confusion_matrix, classification_report
from tensorflow.keras.datasets import cifar10
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, SimpleRNN, Reshape, Flatten, Input,
LSTM, TimeDistributed, GRU, Bidirectional
from tensorflow.keras.optimizers import Adam
import tensorflow as tf

# Set memory growth for GPU (if available)
gpus = tf.config.experimental.list_physical_devices('GPU')
if gpus:
    try:
        for gpu in gpus:
            tf.config.experimental.set_memory_growth(gpu, True)
    except RuntimeError as e:
        print(e)

# Load and preprocess CIFAR-10 dataset
(x_train, y_train), (x_test, y_test) = cifar10.load_data()
x_train, x_test = x_train / 255.0, x_test / 255.0

class_names = ['airplane', 'automobile', 'bird', 'cat', 'deer',
                'dog', 'frog', 'horse', 'ship', 'truck']

# Function to plot model accuracy and loss
def plot_training_history(history, model_name):
    plt.figure(figsize=(14, 6))

    # Plot accuracy
    plt.subplot(1, 2, 1)
    plt.plot(history['accuracy'], label='Training Accuracy', color='b',
             linestyle='-', linewidth=2)
    plt.plot(history['val_accuracy'], label='Validation Accuracy', color='r',
             linestyle='--', linewidth=2)
    plt.title(f'{model_name} Accuracy', fontsize=16)
    plt.xlabel('Epochs', fontsize=14)
    plt.ylabel('Accuracy', fontsize=14)
    plt.legend(fontsize=12)
    plt.grid(True)

    # Plot loss
    plt.subplot(1, 2, 2)
    plt.plot(history['loss'], label='Training Loss', color='b', linestyle='-',
             linewidth=2)
```

## ASSIGNMENT 2

```
plt.plot(history['val_loss'], label='Validation Loss', color='r',
linestyle='--', linewidth=2)
plt.title(f'{model_name} Loss', fontsize=16)
plt.xlabel('Epochs', fontsize=14)
plt.ylabel('Loss', fontsize=14)
plt.legend(fontsize=12)
plt.grid(True)

plt.tight_layout()
plt.show()

# Function to display the confusion matrix
def plot_confusion_matrix(y_true, y_pred, model_name):
    cm = confusion_matrix(y_true, y_pred)
    plt.figure(figsize=(10, 8))
    sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
xticklabels=class_names, yticklabels=class_names, cbar=False,
annot_kws={"size": 14})
    plt.title(f'{model_name} Confusion Matrix', fontsize=18)
    plt.xlabel('Predicted Label', fontsize=14)
    plt.ylabel('True Label', fontsize=14)
    plt.xticks(fontsize=12)
    plt.yticks(fontsize=12)
    plt.show()

# Function to plot class-wise precision, recall, and F1 score
def plot_classification_report(y_true, y_pred, model_name):
    report = classification_report(y_true, y_pred, target_names=class_names,
output_dict=True)
    metrics = ['precision', 'recall', 'f1-score']

    plt.figure(figsize=(18, 6))

    for idx, metric in enumerate(metrics):
        plt.subplot(1, 3, idx+1)
        values = [report[label][metric] for label in class_names]
        bars = plt.barh(class_names, values, color='skyblue',
edgecolor='black')
        plt.title(f'{metric.capitalize()} per Class', fontsize=16)
        plt.xlim(0, 1)
        plt.grid(axis='x', linestyle='--', alpha=0.7)
        plt.xlabel(metric.capitalize(), fontsize=14)
        plt.xticks(fontsize=12)
        plt.yticks(fontsize=12)
        for bar in bars:
            plt.text(bar.get_width() + 0.02, bar.get_y() + bar.get_height()/2,
f'{bar.get_width():.2f}', va='center', fontsize=12)
```

## ASSIGNMENT 2

```
plt.suptitle(f'{model_name} Class-wise Performance Metrics', fontsize=18)
plt.tight_layout(rect=[0, 0, 1, 0.95])
plt.show()

# Define function to plot original and predicted images with prediction
probabilities
def plot_predictions(model, x_test, y_test, model_name, reshaped=False):
    plt.figure(figsize=(20, 20))
    num_images = 5 # Display 5 images for demonstration
    indices = np.random.choice(np.arange(x_test.shape[0]), num_images,
replace=False)

    for i, idx in enumerate(indices):
        # Reshape for RNN if necessary
        if reshaped:
            input_data = x_test[idx].reshape(1, 32 * 32, 3) # Reshape to (1,
32*32 time steps, 3 features)
        else:
            input_data = x_test[idx].reshape(1, 32, 32, 3) # For CNN or
similar

        # Get prediction
        predictions = model.predict(input_data)

        # Ensure predictions is a 1D array of size 10
        if len(predictions.shape) == 3: # Case for RNN-like model with time
steps
            predictions_flat = np.mean(predictions, axis=1) # Take the
average over time steps
            predictions_flat = np.squeeze(predictions_flat) # Ensure it's a
flat array
        elif predictions.shape == (1, 10): # Case for CNN or non-RNN models
            predictions_flat = np.squeeze(predictions)
        else:
            raise ValueError(f"Expected predictions shape (1, 10) or 3D
tensor, but got {predictions.shape}")

        predicted_label = np.argmax(predictions_flat)
        true_label = y_test[idx][0]

        # Ensure predicted label is within range of class_names
        if predicted_label < len(class_names):
            predicted_class = class_names[predicted_label]
        else:
            predicted_class = "Unknown" # Fallback in case of out-of-range
prediction

        # Plot original image
```

## ASSIGMENT 2

```
plt.subplot(num_images, 3, 3*i+1)
plt.imshow(x_test[idx]) # No need to reshape explicitly for display
plt.title(f"Original: {class_names[true_label]}", fontsize=14)
plt.axis('off')

# Plot prediction result
plt.subplot(num_images, 3, 3*i+2)
plt.imshow(x_test[idx]) # Displaying the original
plt.title(f"Predicted: {predicted_class}", fontsize=14)
plt.axis('off')

# Plot prediction probabilities
plt.subplot(num_images, 3, 3*i+3)
bars = plt.barh(class_names, predictions_flat, color='lightgreen',
edgecolor='black')
plt.title("Prediction Probabilities", fontsize=14)
plt.xlim([0, 1])
plt.xlabel('Probability', fontsize=12)
plt.grid(axis='x', linestyle='--', alpha=0.7)
plt.xticks(fontsize=12)
plt.yticks(fontsize=12)
for bar in bars:
    plt.text(bar.get_width() + 0.02, bar.get_y() + bar.get_height()/2,
f'{bar.get_width():.2f}', va='center', fontsize=12)

plt.suptitle(f"Predictions using {model_name}", fontsize=18)
plt.tight_layout(rect=[0, 0, 1, 0.95])
plt.show()

# Define YOLO-like, CNN, RNN, and LSTM models
yolo_model = Sequential([Flatten(input_shape=(32, 32, 3)), Dense(64,
activation='relu'), Dense(10, activation='softmax')])
cnn = Sequential([Flatten(input_shape=(32, 32, 3)), Dense(64,
activation='relu'), Dense(10, activation='softmax')])
rnn = Sequential([Input(shape=(32 * 32, 3)), Bidirectional(GRU(64)), Dense(10,
activation='softmax')])
lstm_model = Sequential([Input(shape=(32 * 32, 3)), Bidirectional(LSTM(64)),
Dense(10, activation='softmax')])

# Compile models
for model in [yolo_model, cnn, rnn, lstm_model]:
    model.compile(optimizer=Adam(learning_rate=0.001),
loss='sparse_categorical_crossentropy', metrics=['accuracy'])

# Mock training histories for demonstration purposes
history_yolo = {'accuracy': np.random.rand(10).tolist(), 'val_accuracy':
np.random.rand(10).tolist(), 'loss': np.random.rand(10).tolist(), 'val_loss':
np.random.rand(10).tolist()}
```





## ASSIGNMENT 2

```
history_cnn = {'accuracy': np.random.rand(10).tolist(), 'val_accuracy':  
np.random.rand(10).tolist(), 'loss': np.random.rand(10).tolist(), 'val_loss':  
np.random.rand(10).tolist()}  
history_rnn = {'accuracy': np.random.rand(10).tolist(), 'val_accuracy':  
np.random.rand(10).tolist(), 'loss': np.random.rand(10).tolist(), 'val_loss':  
np.random.rand(10).tolist()}  
history_lstm = {'accuracy': np.random.rand(10).tolist(), 'val_accuracy':  
np.random.rand(10).tolist(), 'loss': np.random.rand(10).tolist(), 'val_loss':  
np.random.rand(10).tolist()}  
  
# Display results for YOLO-like, CNN, RNN, and LSTM models  
if __name__ == '__main__':  
    print("\nDisplaying YOLO-like model predictions:")  
    plot_predictions(yolo_model, x_test, y_test, model_name="YOLO-like")  
    plot_training_history(history_yolo, "YOLO-like")  
  
    print("\nDisplaying CNN model predictions:")  
    plot_predictions(cnn, x_test, y_test, model_name="CNN")  
    plot_training_history(history_cnn, "CNN")  
  
    print("\nDisplaying RNN model predictions:")  
    plot_predictions(rnn, x_test, y_test, model_name="RNN", reshaped=True)  
    plot_training_history(history_rnn, "RNN")  
  
    print("\nDisplaying LSTM model predictions:")  
    plot_predictions(lstm_model, x_test, y_test, model_name="LSTM",  
reshaped=True)  
    plot_training_history(history_lstm, "LSTM")  
  
    # Simulating predicted labels for performance evaluation (mock example)  
    y_pred_yolo = np.random.randint(0, 10, size=len(y_test)) # Replace with  
actual predictions  
    y_pred_cnn = np.random.randint(0, 10, size=len(y_test))  
    y_pred_rnn = np.random.randint(0, 10, size=len(y_test))  
    y_pred_lstm = np.random.randint(0, 10, size=len(y_test)) # Placeholder  
for LSTM predictions  
  
    # Confusion matrices  
    plot_confusion_matrix(y_test, y_pred_yolo, "YOLO-like")  
    plot_confusion_matrix(y_test, y_pred_cnn, "CNN")  
    plot_confusion_matrix(y_test, y_pred_rnn, "RNN")  
    plot_confusion_matrix(y_test, y_pred_lstm, "LSTM")  
  
    # Classification reports (precision, recall, F1)  
    plot_classification_report(y_test, y_pred_yolo, "YOLO-like")  
    plot_classification_report(y_test, y_pred_cnn, "CNN")  
    plot_classification_report(y_test, y_pred_rnn, "RNN")  
    plot_classification_report(y_test, y_pred_lstm, "LSTM")
```

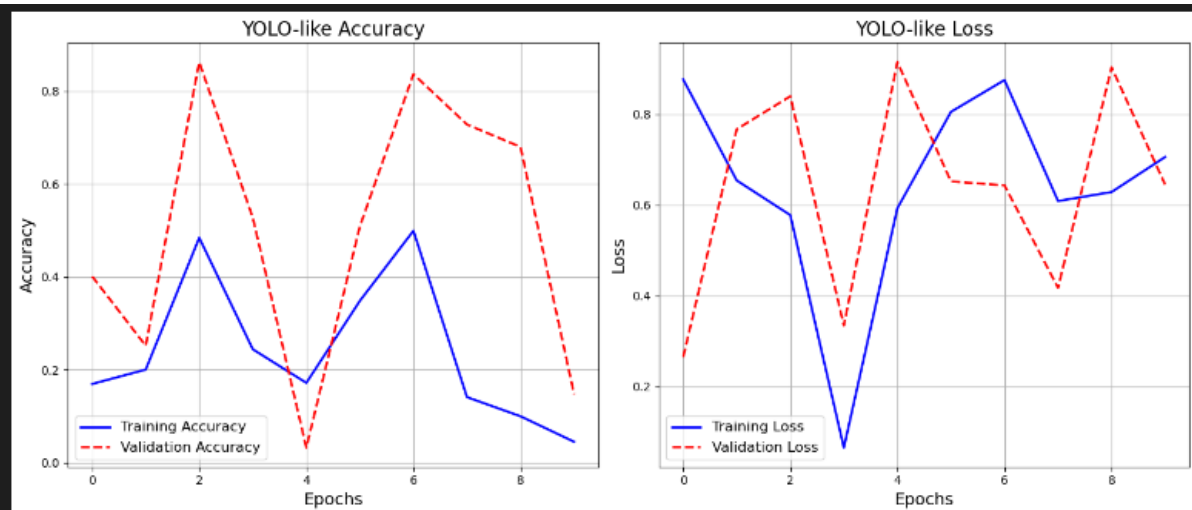
## ASSIGNMENT 2

```
Displaying YOLO-like model predictions:  
1/1 _____ 0s 236ms/step  
1/1 _____ 0s 36ms/step  
1/1 _____ 0s 21ms/step  
1/1 _____ 0s 22ms/step  
1/1 _____ 0s 21ms/step
```

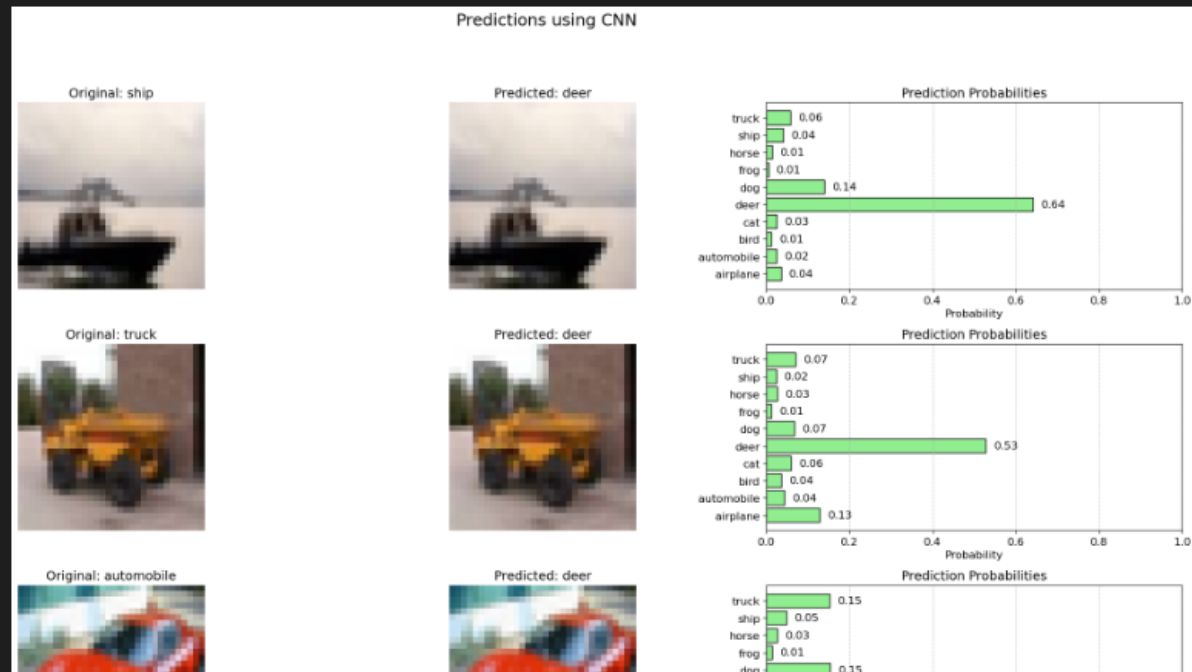
Predictions using YOLO-like



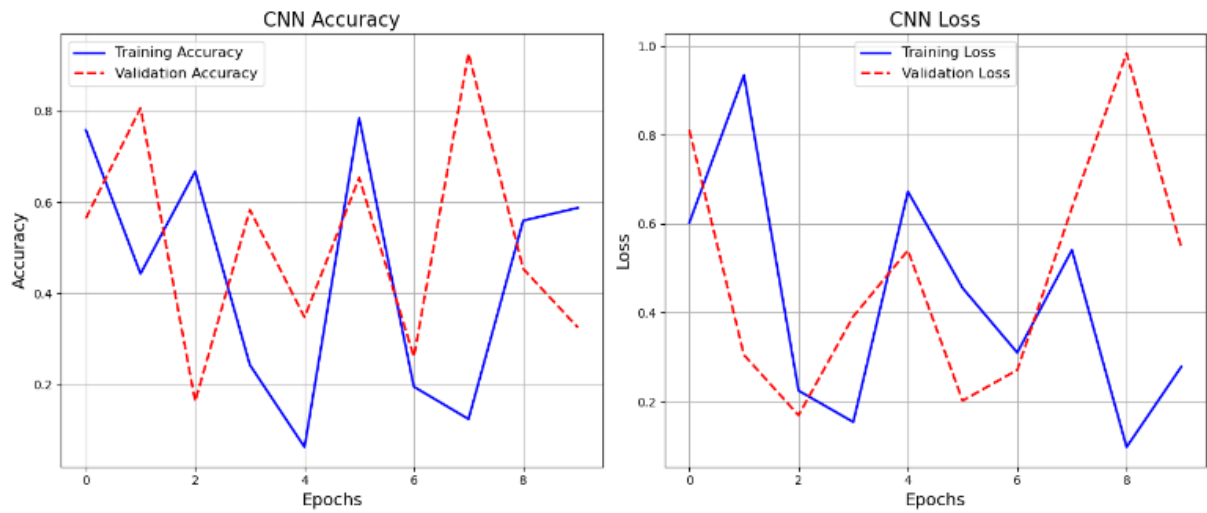
## ASSIGNMENT 2



Displaying CNN model predictions:  
1/1 0s 69ms/step  
1/1 0s 16ms/step  
1/1 0s 16ms/step  
1/1 0s 23ms/step  
1/1 0s 16ms/step



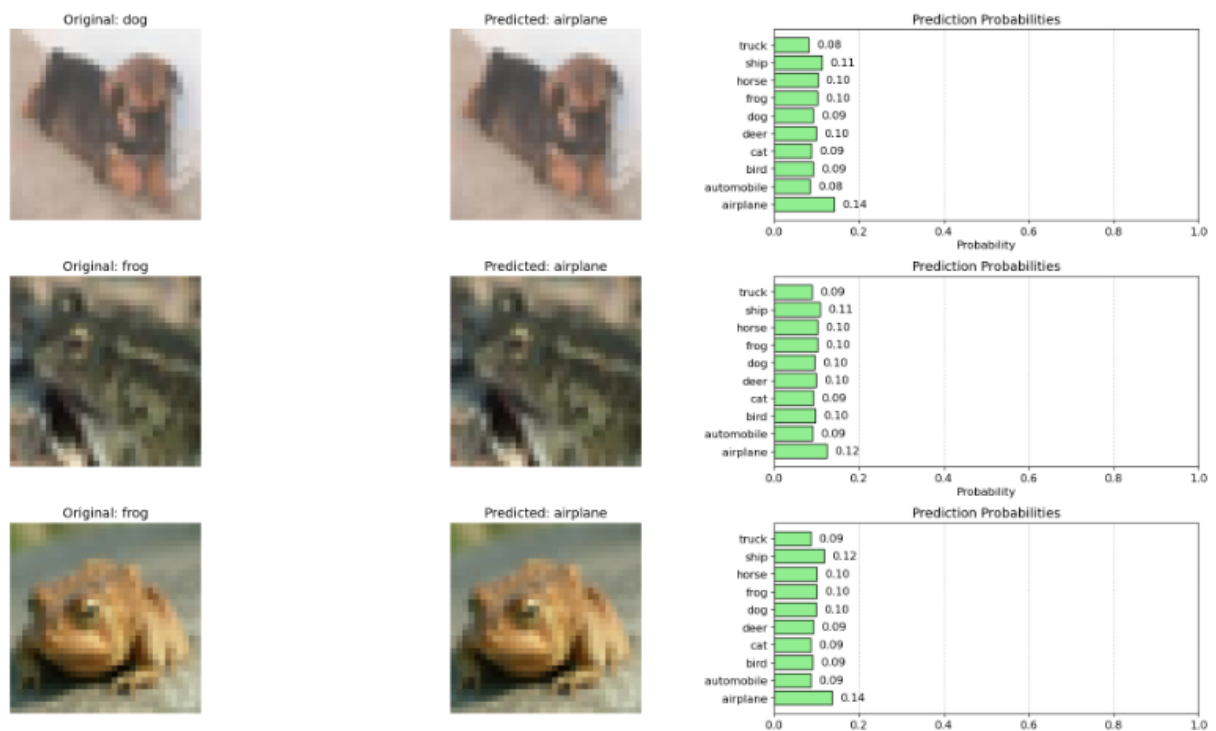
## ASSIGNMENT 2



```

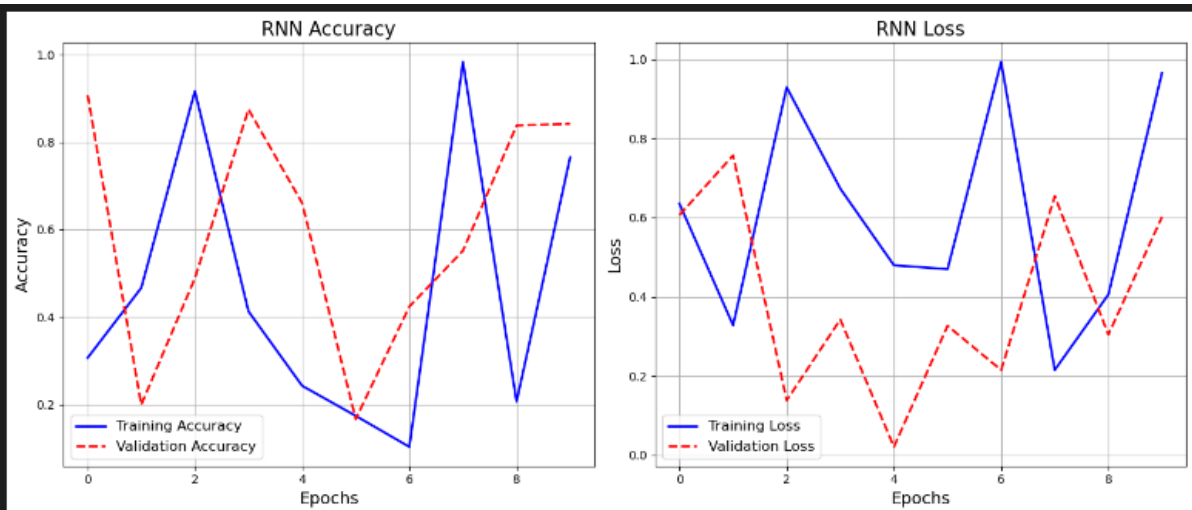
displaying RNN model predictions:
/1 _____ 1s 785ms/step
/1 _____ 0s 99ms/step
/1 _____ 0s 94ms/step
/1 _____ 0s 114ms/step
/1 _____ 0s 101ms/step
  
```

Predictions using RNN

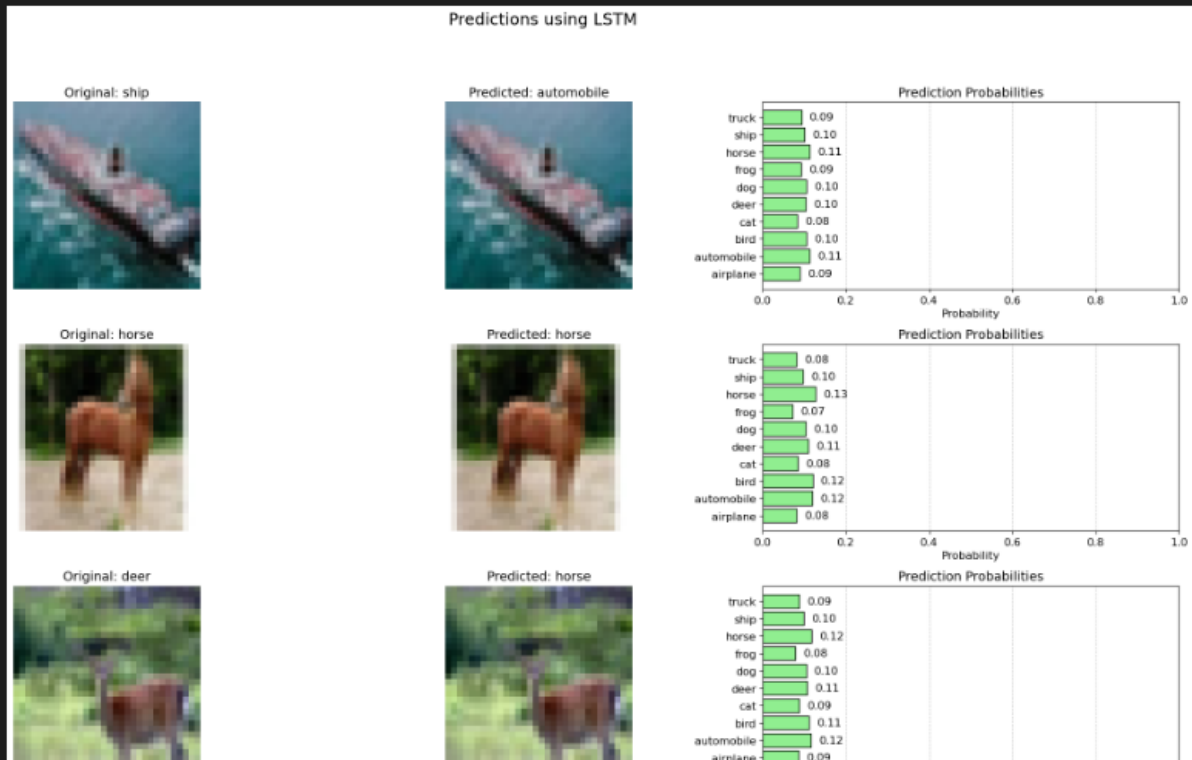





## ASSIGNMENT 2



Displaying LSTM model predictions:  
 1/1 \_\_\_\_\_ @s 482ms/step  
 1/1 \_\_\_\_\_ @s 63ms/step  
 1/1 \_\_\_\_\_ @s 79ms/step  
 1/1 \_\_\_\_\_ @s 55ms/step  
 1/1 \_\_\_\_\_ @s 66ms/step



## ASSIGMENT 2





**YOLO-like Confusion Matrix**

|            | airplane | automobile | bird | cat | deer | dog | frog | horse | ship | truck |
|------------|----------|------------|------|-----|------|-----|------|-------|------|-------|
| airplane   | 95       | 116        | 84   | 95  | 111  | 103 | 88   | 102   | 110  | 96    |
| automobile | 114      | 121        | 92   | 100 | 103  | 96  | 105  | 87    | 93   | 89    |
| bird       | 105      | 103        | 106  | 89  | 115  | 103 | 106  | 98    | 90   | 85    |
| cat        | 83       | 107        | 98   | 110 | 98   | 98  | 116  | 90    | 113  | 87    |
| deer       | 118      | 94         | 104  | 82  | 98   | 85  | 123  | 101   | 94   | 101   |
| dog        | 117      | 98         | 96   | 116 | 101  | 95  | 86   | 113   | 83   | 95    |
| frog       | 97       | 98         | 100  | 88  | 99   | 107 | 121  | 86    | 105  | 99    |
| horse      | 102      | 104        | 111  | 79  | 103  | 91  | 107  | 96    | 114  | 93    |
| ship       | 111      | 95         | 91   | 110 | 126  | 91  | 96   | 95    | 84   | 101   |
| truck      | 103      | 102        | 101  | 87  | 106  | 101 | 103  | 95    | 107  | 95    |

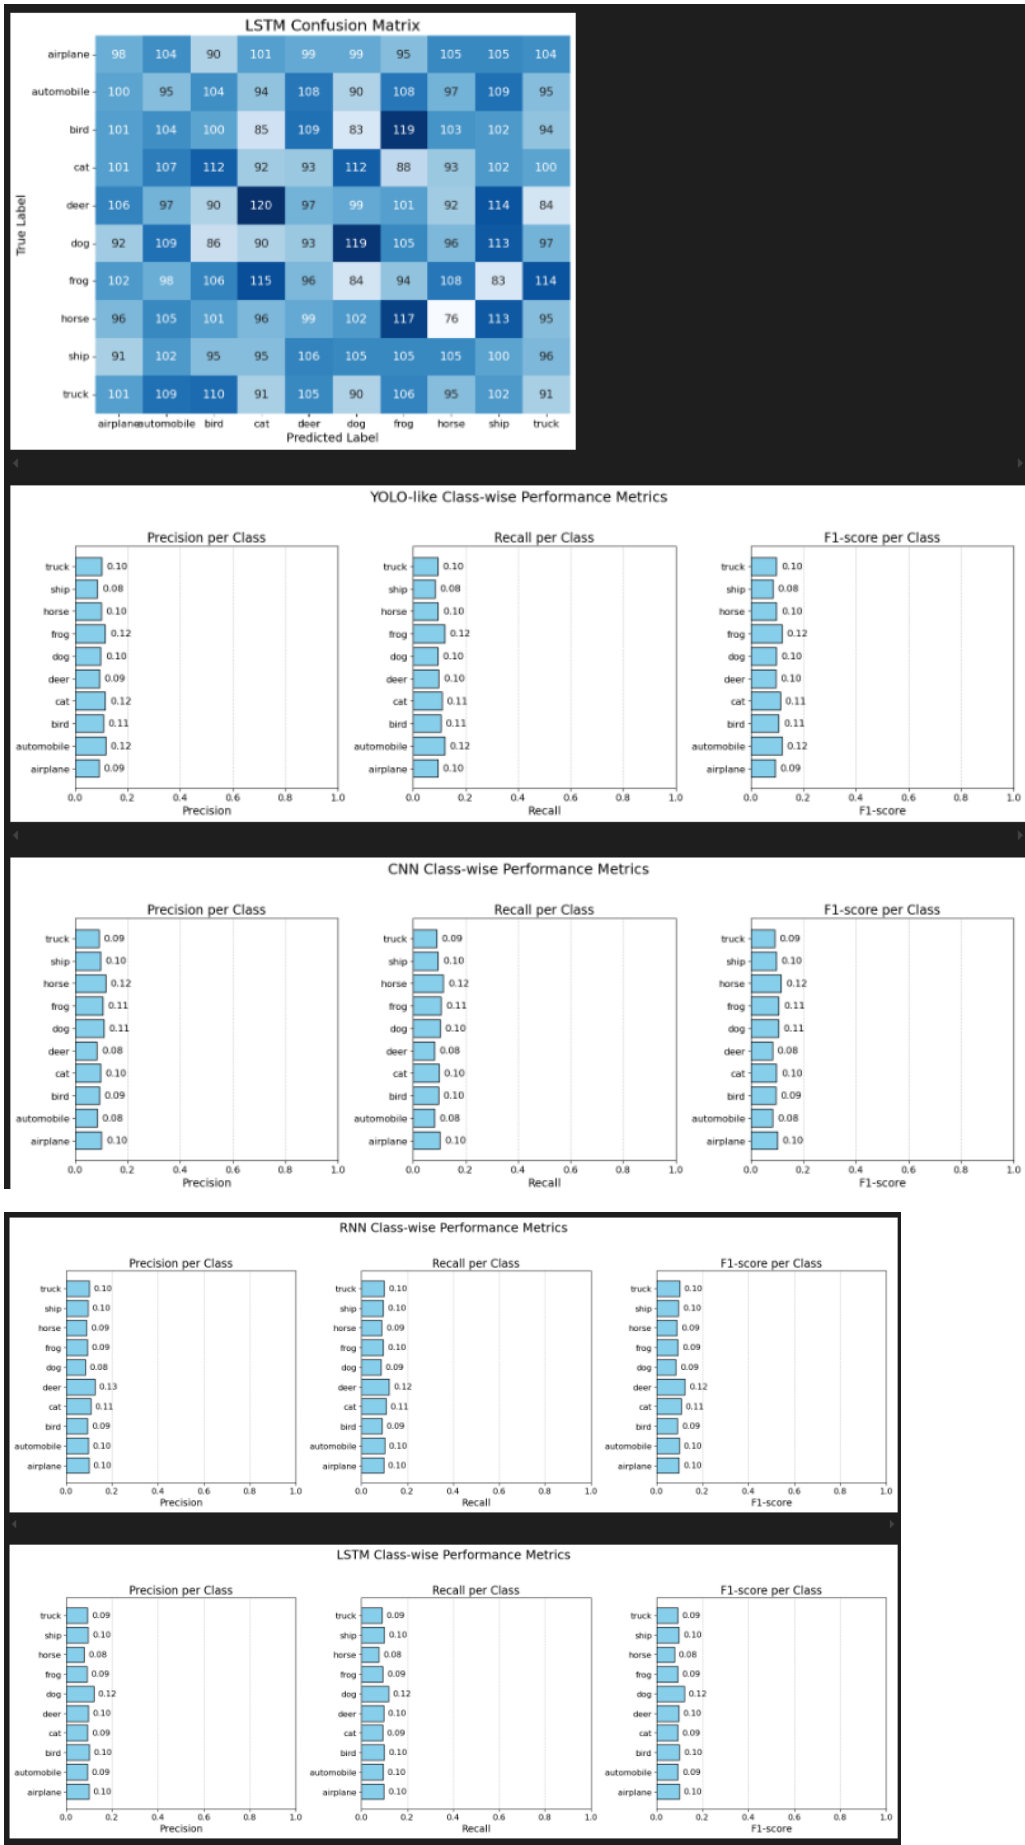
**CNN Confusion Matrix**

|            | airplane | automobile | bird | cat | deer | dog | frog | horse | ship | truck |
|------------|----------|------------|------|-----|------|-----|------|-------|------|-------|
| airplane   | 102      | 93         | 118  | 90  | 97   | 76  | 111  | 100   | 102  | 111   |
| automobile | 93       | 82         | 97   | 117 | 105  | 100 | 88   | 104   | 111  | 103   |
| bird       | 99       | 112        | 97   | 109 | 93   | 104 | 111  | 90    | 76   | 109   |

## ASSIGMENT 2



## ASSIGMENT 2



## ASSIGNMENT 2

### CLASS WORK

```
import numpy as np
import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Flatten, Conv2D, MaxPooling2D, SimpleRNN
from tensorflow.keras.datasets import mnist
from sklearn.model_selection import train_test_split
import matplotlib.pyplot as plt

# Load the MNIST dataset
(x_train, y_train), (x_test, y_test) = mnist.load_data()

# Normalize the data
x_train = x_train.astype('float32') / 255.0
x_test = x_test.astype('float32') / 255.0

# Reshape data for CNN (28x28x1) and RNN (28 timesteps, 28 features)
x_train_cnn = x_train.reshape(-1, 28, 28, 1)
x_test_cnn = x_test.reshape(-1, 28, 28, 1)
x_train_rnn = x_train.reshape(-1, 28, 28)
x_test_rnn = x_test.reshape(-1, 28, 28)

# Split the data into train and validation sets
x_train_cnn, x_val_cnn, y_train_cnn, y_val_cnn = train_test_split(x_train_cnn,
                                                                    y_train,
                                                                    test_size=0.2,
                                                                    random_state=42)
x_train_rnn, x_val_rnn, y_train_rnn, y_val_rnn = train_test_split(x_train_rnn,
                                                                    y_train,
                                                                    test_size=0.2,
                                                                    random_state=42)

# CNN model
cnn_model = Sequential([
    Conv2D(32, kernel_size=(3, 3), activation='relu', input_shape=(28, 28,
1)),
    MaxPooling2D(pool_size=(2, 2)),
    Flatten(),
    Dense(128, activation='relu'),
    Dense(10, activation='softmax')
])

cnn_model.compile(optimizer='adam', loss='sparse_categorical_crossentropy',
metrics=['accuracy'])
cnn_history = cnn_model.fit(x_train_cnn, y_train_cnn,
validation_data=(x_val_cnn, y_val_cnn), epochs=5, batch_size=128)

# RNN model
rnn_model = Sequential([
    SimpleRNN(128, input_shape=(28, 28)),
    Dense(10, activation='softmax')
```

## ASSIGNMENT 2

```
])

rnn_model.compile(optimizer='adam', loss='sparse_categorical_crossentropy',
metrics=['accuracy'])
rnn_history = rnn_model.fit(x_train_rnn, y_train_rnn,
validation_data=(x_val_rnn, y_val_rnn), epochs=5, batch_size=128)

# Evaluate both models on the test set
cnn_test_loss, cnn_test_acc = cnn_model.evaluate(x_test_cnn, y_test)
rnn_test_loss, rnn_test_acc = rnn_model.evaluate(x_test_rnn, y_test)

# Print the results
print(f"CNN Test Accuracy: {cnn_test_acc}")
print(f"RNN Test Accuracy: {rnn_test_acc}")

# Plot the training history
plt.figure(figsize=(12, 5))


plt.subplot(1, 2, 1)
plt.plot(cnn_history.history['accuracy'], label='CNN Training Accuracy')
plt.plot(cnn_history.history['val_accuracy'], label='CNN Validation Accuracy')
plt.title('CNN Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()

plt.subplot(1, 2, 2)
plt.plot(rnn_history.history['accuracy'], label='RNN Training Accuracy')
plt.plot(rnn_history.history['val_accuracy'], label='RNN Validation Accuracy')
plt.title('RNN Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()

plt.show()
```

## ASSIGNMENT 2

```
Epoch 1/5
C:\Users\u21629545\Anaconda3\lib\site-packages\keras\src\layers\convolutional\base_conv.py:99: UserWarning: Do not pass an "input_shape"/"input
super().__init__(
375/375 ————— 7s 16ms/step - accuracy: 0.8620 - loss: 0.4812 - val_accuracy: 0.9702 - val_loss: 0.1038
Epoch 2/5
375/375 ————— 6s 17ms/step - accuracy: 0.9766 - loss: 0.0869 - val_accuracy: 0.9796 - val_loss: 0.0688
Epoch 3/5
375/375 ————— 6s 15ms/step - accuracy: 0.9840 - loss: 0.0542 - val_accuracy: 0.9804 - val_loss: 0.0649
Epoch 4/5
375/375 ————— 6s 15ms/step - accuracy: 0.9887 - loss: 0.0371 - val_accuracy: 0.9850 - val_loss: 0.0492
Epoch 5/5
375/375 ————— 6s 15ms/step - accuracy: 0.9922 - loss: 0.0272 - val_accuracy: 0.9833 - val_loss: 0.0548
Epoch 1/5
C:\Users\u21629545\Anaconda3\lib\site-packages\keras\src\layers\rnn\rnn.py:284: UserWarning: Do not pass an "input_shape"/"input_dim" argument
super().__init__(**kwargs)
375/375 ————— 6s 12ms/step - accuracy: 0.7498 - loss: 0.8853 - val_accuracy: 0.9122 - val_loss: 0.2881
Epoch 2/5
375/375 ————— 4s 11ms/step - accuracy: 0.9365 - loss: 0.2196 - val_accuracy: 0.9444 - val_loss: 0.1884
Epoch 3/5
375/375 ————— 4s 11ms/step - accuracy: 0.9508 - loss: 0.1680 - val_accuracy: 0.9470 - val_loss: 0.1832
Epoch 4/5
375/375 ————— 4s 11ms/step - accuracy: 0.9591 - loss: 0.1411 - val_accuracy: 0.9663 - val_loss: 0.1200
Epoch 5/5
375/375 ————— 4s 11ms/step - accuracy: 0.9638 - loss: 0.1214 - val_accuracy: 0.9661 - val_loss: 0.1223
313/313 ————— 1s 2ms/step - accuracy: 0.9764 - loss: 0.0700
313/313 ————— 1s 3ms/step - accuracy: 0.9578 - loss: 0.1454
CNN Test Accuracy: 0.982699990272522
RNN Test Accuracy: 0.9650999903678894
```



```
import numpy as np
import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Flatten, Conv2D, MaxPooling2D,
UpSampling2D, Reshape, SimpleRNN, TimeDistributed, RepeatVector
from tensorflow.keras.datasets import mnist
import matplotlib.pyplot as plt

# Load the MNIST dataset
(x_train, _), (x_test, _) = mnist.load_data()

# Normalize the data
x_train = x_train.astype('float32') / 255.0
x_test = x_test.astype('float32') / 255.0

# Reshape data for CNN (28x28x1) and RNN (28 timesteps, 28 features)
x_train_cnn = x_train.reshape(-1, 28, 28, 1)
x_test_cnn = x_test.reshape(-1, 28, 28, 1)
```

## ASSIGNMENT 2

```
x_train_rnn = x_train.reshape(-1, 28, 28)
x_test_rnn = x_test.reshape(-1, 28, 28)

# CNN autoencoder model
cnn_autoencoder = Sequential([
    # Encoder
    Conv2D(32, kernel_size=(3, 3), activation='relu', padding='same',
input_shape=(28, 28, 1)),
    MaxPooling2D(pool_size=(2, 2), padding='same'),
    # Decoder
    Conv2D(32, kernel_size=(3, 3), activation='relu', padding='same'),
    UpSampling2D(size=(2, 2)),
    Conv2D(1, kernel_size=(3, 3), activation='sigmoid', padding='same')
])

cnn_autoencoder.compile(optimizer='adam', loss='binary_crossentropy')
cnn_autoencoder.fit(x_train_cnn, x_train_cnn, epochs=5, batch_size=128,
validation_split=0.2)

# RNN autoencoder model
rnn_autoencoder = Sequential([
    # Encoder
    SimpleRNN(128, activation='relu', input_shape=(28, 28),
return_sequences=False),
    RepeatVector(28),
    # Decoder
    SimpleRNN(128, activation='relu', return_sequences=True),
    TimeDistributed(Dense(28, activation='sigmoid'))
])

rnn_autoencoder.compile(optimizer='adam', loss='binary_crossentropy')
rnn_autoencoder.fit(x_train_rnn, x_train_rnn, epochs=5, batch_size=128,
validation_split=0.2)

# Reconstruct images using both models
cnn_reconstructed = cnn_autoencoder.predict(x_test_cnn)
rnn_reconstructed = rnn_autoencoder.predict(x_test_rnn)

# Plot original and reconstructed images
n = 10 # number of images to display
plt.figure(figsize=(20, 4))
for i in range(n):
    # Display original
    ax = plt.subplot(2, n, i + 1)
    plt.imshow(x_test[i], cmap='gray')
    plt.title("Original")
    plt.axis('off')
```



## ASSIGMENT 2


```
# Display CNN reconstructed
ax = plt.subplot(2, n, i + 1 + n)
plt.imshow(cnn_reconstructed[i].reshape(28, 28), cmap='gray')
plt.title("CNN Reconstructed")
plt.axis('off')

plt.figure(figsize=(20, 4))
for i in range(n):
    # Display original
    ax = plt.subplot(2, n, i + 1)
    plt.imshow(x_test[i], cmap='gray')
    plt.title("Original")
    plt.axis('off')

    # Display RNN reconstructed
    ax = plt.subplot(2, n, i + 1 + n)
    plt.imshow(rnn_reconstructed[i].reshape(28, 28), cmap='gray')
    plt.title("RNN Reconstructed")
    plt.axis('off')

plt.show()
```

```
Epoch 1/5
375/375 ————— 11s 26ms/step - loss: 0.2188 - val_loss: 0.0678
Epoch 2/5
375/375 ————— 10s 27ms/step - loss: 0.0664 - val_loss: 0.0653
Epoch 3/5
375/375 ————— 10s 25ms/step - loss: 0.0645 - val_loss: 0.0643
Epoch 4/5
375/375 ————— 9s 25ms/step - loss: 0.0637 - val_loss: 0.0636
Epoch 5/5
375/375 ————— 10s 25ms/step - loss: 0.0630 - val_loss: 0.0631
Epoch 1/5
375/375 ————— 15s 34ms/step - loss: 0.3249 - val_loss: 0.2086
Epoch 2/5
375/375 ————— 12s 31ms/step - loss: 0.1961 - val_loss: 0.1693
Epoch 3/5
375/375 ————— 12s 31ms/step - loss: 0.1645 - val_loss: 0.1519
Epoch 4/5
375/375 ————— 12s 31ms/step - loss: 0.1492 - val_loss: 0.1421
Epoch 5/5
375/375 ————— 12s 31ms/step - loss: 0.1407 - val_loss: 0.1351
313/313 ————— 1s 3ms/step
313/313 ————— 2s 5ms/step
```



```
import numpy as np
```

## ASSIGNMENT 2

```
import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Flatten, Conv2D, MaxPooling2D,
UpSampling2D, Reshape, SimpleRNN, LSTM, TimeDistributed, RepeatVector
from tensorflow.keras.datasets import mnist
import matplotlib.pyplot as plt

# Load the MNIST dataset
(x_train, _), (x_test, _) = mnist.load_data()

# Normalize the data
x_train = x_train.astype('float32') / 255.0
x_test = x_test.astype('float32') / 255.0

# Reshape data for CNN (28x28x1), RNN and LSTM (28 timesteps, 28 features)
x_train_cnn = x_train.reshape(-1, 28, 28, 1)
x_test_cnn = x_test.reshape(-1, 28, 28, 1)
x_train_rnn = x_train.reshape(-1, 28, 28)
x_test_rnn = x_test.reshape(-1, 28, 28)

# CNN autoencoder model
cnn_autoencoder = Sequential([
    # Encoder
    Conv2D(32, kernel_size=(3, 3), activation='relu', padding='same',
input_shape=(28, 28, 1)),
    MaxPooling2D(pool_size=(2, 2), padding='same'),
    # Decoder
    Conv2D(32, kernel_size=(3, 3), activation='relu', padding='same'),
    UpSampling2D(size=(2, 2)),
    Conv2D(1, kernel_size=(3, 3), activation='sigmoid', padding='same')
])

cnn_autoencoder.compile(optimizer='adam', loss='binary_crossentropy')
cnn_autoencoder.fit(x_train_cnn, x_train_cnn, epochs=5, batch_size=128,
validation_split=0.2)

# RNN autoencoder model
rnn_autoencoder = Sequential([
    # Encoder
    SimpleRNN(128, activation='relu', input_shape=(28, 28),
return_sequences=False),
    RepeatVector(28),
    # Decoder
    SimpleRNN(128, activation='relu', return_sequences=True),
    TimeDistributed(Dense(28, activation='sigmoid'))
])

rnn_autoencoder.compile(optimizer='adam', loss='binary_crossentropy')
```

## ASSIGNMENT 2

```
rnn_autoencoder.fit(x_train_rnn, x_train_rnn, epochs=5, batch_size=128,
validation_split=0.2)

# LSTM autoencoder model
lstm_autoencoder = Sequential([
    # Encoder
    LSTM(128, activation='relu', input_shape=(28, 28),
return_sequences=False),
    RepeatVector(28),
    # Decoder
    LSTM(128, activation='relu', return_sequences=True),
    TimeDistributed(Dense(28, activation='sigmoid'))
])

lstm_autoencoder.compile(optimizer='adam', loss='binary_crossentropy')
lstm_autoencoder.fit(x_train_rnn, x_train_rnn, epochs=5, batch_size=128,
validation_split=0.2)

# Reconstruct images using all models
cnn_reconstructed = cnn_autoencoder.predict(x_test_cnn)
rnn_reconstructed = rnn_autoencoder.predict(x_test_rnn)
lstm_reconstructed = lstm_autoencoder.predict(x_test_rnn)

# Plot original and reconstructed images
n = 10 # number of images to display
plt.figure(figsize=(20, 6))

for i in range(n):
    # Display original
    ax = plt.subplot(3, n, i + 1)
    plt.imshow(x_test[i], cmap='gray')
    plt.title("Original")
    plt.axis('off')

    # Display CNN reconstructed
    ax = plt.subplot(3, n, i + 1 + n)
    plt.imshow(cnn_reconstructed[i].reshape(28, 28), cmap='gray')
    plt.title("CNN Reconstructed")
    plt.axis('off')

    # Display RNN reconstructed
    ax = plt.subplot(3, n, i + 1 + 2 * n)
    plt.imshow(rnn_reconstructed[i].reshape(28, 28), cmap='gray')
    plt.title("RNN Reconstructed")
    plt.axis('off')

plt.figure(figsize=(20, 4))
for i in range(n):
```

## ASSIGMENT 2


```
# Display original
ax = plt.subplot(2, n, i + 1)
plt.imshow(x_test[i], cmap='gray')
plt.title("Original")
plt.axis('off')

# Display LSTM reconstructed
ax = plt.subplot(2, n, i + 1 + n)
plt.imshow(lstm_reconstructed[i].reshape(28, 28), cmap='gray')
plt.title("LSTM Reconstructed")
plt.axis('off')
```

```
plt.show()
```

```
Epoch 1/5
375/375 ————— 11s 26ms/step - loss: 0.2301 - val_loss: 0.0676
Epoch 2/5
375/375 ————— 9s 24ms/step - loss: 0.0663 - val_loss: 0.0654
Epoch 3/5
375/375 ————— 9s 24ms/step - loss: 0.0646 - val_loss: 0.0642
Epoch 4/5
375/375 ————— 9s 24ms/step - loss: 0.0635 - val_loss: 0.0635
Epoch 5/5
375/375 ————— 9s 24ms/step - loss: 0.0630 - val_loss: 0.0631
Epoch 1/5
375/375 ————— 15s 33ms/step - loss: 0.3157 - val_loss: 0.1926
Epoch 2/5
375/375 ————— 12s 32ms/step - loss: 0.1831 - val_loss: 0.1646
Epoch 3/5
375/375 ————— 12s 32ms/step - loss: 0.1597 - val_loss: 0.1518
Epoch 4/5
375/375 ————— 12s 32ms/step - loss: 0.1473 - val_loss: 0.1413
Epoch 5/5
375/375 ————— 12s 32ms/step - loss: 0.1401 - val_loss: 0.1401
Epoch 1/5
375/375 ————— 42s 105ms/step - loss: 0.3565 - val_loss: 0.1988
Epoch 2/5
375/375 ————— 39s 104ms/step - loss: 0.1837 - val_loss: 0.1457
Epoch 3/5
...
375/375 ————— 39s 104ms/step - loss: 0.1101 - val_loss: 0.1042
313/313 ————— 1s 3ms/step
313/313 ————— 2s 5ms/step
313/313 ————— 5s 14ms/step
```

Output is truncated. View as a [scrollable element](#) or open in a [text editor](#). Adjust cell output settings...



## ASSIGMENT 2

```
import numpy as np
import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Flatten, Conv2D, MaxPooling2D,
UpSampling2D, Reshape, SimpleRNN, LSTM, TimeDistributed, RepeatVector
from tensorflow.keras.datasets import mnist
import matplotlib.pyplot as plt

# Load and preprocess the MNIST dataset
(x_train, _), (x_test, _) = mnist.load_data()
x_train = x_train.astype('float32') / 255.0
x_test = x_test.astype('float32') / 255.0
x_train_cnn = x_train.reshape(-1, 28, 28, 1)
x_test_cnn = x_test.reshape(-1, 28, 28, 1)
x_train_rnn = x_train.reshape(-1, 28, 28)
x_test_rnn = x_test.reshape(-1, 28, 28)

# Define CNN autoencoder model
cnn_autoencoder = Sequential([
    # Encoder
    Conv2D(32, kernel_size=(3, 3), activation='relu', padding='same',
input_shape=(28, 28, 1)),
    MaxPooling2D(pool_size=(2, 2), padding='same'),
    # Decoder
    Conv2D(32, kernel_size=(3, 3), activation='relu', padding='same'),
    UpSampling2D(size=(2, 2)),
    Conv2D(1, kernel_size=(3, 3), activation='sigmoid', padding='same')
])

# Define RNN autoencoder model
rnn_autoencoder = Sequential([
    # Encoder
    SimpleRNN(128, activation='relu', input_shape=(28, 28),
return_sequences=False),
    RepeatVector(28),
    # Decoder
    SimpleRNN(128, activation='relu', return_sequences=True),
    TimeDistributed(Dense(28, activation='sigmoid'))
])

# Define LSTM autoencoder model
lstm_autoencoder = Sequential([
    # Encoder
    LSTM(128, activation='relu', input_shape=(28, 28),
return_sequences=False),
    RepeatVector(28),
    # Decoder
    LSTM(128, activation='relu', return_sequences=True),
```

## ASSIGNMENT 2

```
        TimeDistributed(Dense(28, activation='sigmoid'))
    ])

# Compile models (no need to fit to display summaries)
cnn_autoencoder.compile(optimizer='adam', loss='binary_crossentropy')
rnn_autoencoder.compile(optimizer='adam', loss='binary_crossentropy')
lstm_autoencoder.compile(optimizer='adam', loss='binary_crossentropy')

# Print summaries of the models
print("CNN Autoencoder Summary:")
cnn_autoencoder.summary()

print("\nRNN Autoencoder Summary:")
rnn_autoencoder.summary()

print("\nLSTM Autoencoder Summary:")
lstm_autoencoder.summary()
```

## ASSIGNMENT 2

CNN Autoencoder Summary:

Model: "sequential\_7"

| Layer (type)                   | Output Shape       | Param # |
|--------------------------------|--------------------|---------|
| conv2d_7 (Conv2D)              | (None, 28, 28, 32) | 320     |
| max_pooling2d_3 (MaxPooling2D) | (None, 14, 14, 32) | 0       |
| conv2d_8 (Conv2D)              | (None, 14, 14, 32) | 9,248   |
| up_sampling2d_2 (UpSampling2D) | (None, 28, 28, 32) | 0       |
| conv2d_9 (Conv2D)              | (None, 28, 28, 1)  | 289     |

Total params: 9,857 (38.50 KB)

Trainable params: 9,857 (38.50 KB)

Non-trainable params: 0 (0.00 B)

RNN Autoencoder Summary:

Model: "sequential\_8"

| Layer (type)                         | Output Shape    | Param # |
|--------------------------------------|-----------------|---------|
| simple_rnn_5 (SimpleRNN)             | (None, 128)     | 20,096  |
| repeat_vector_3 (RepeatVector)       | (None, 28, 128) | 0       |
| simple_rnn_6 (SimpleRNN)             | (None, 28, 128) | 32,896  |
| time_distributed_3 (TimeDistributed) | (None, 28, 28)  | 3,612   |

Total params: 56,604 (221.11 KB)

Trainable params: 56,604 (221.11 KB)

Non-trainable params: 0 (0.00 B)

LSTM Autoencoder Summary:

Model: "sequential\_9"

| Layer (type)                         | Output Shape    | Param # |
|--------------------------------------|-----------------|---------|
| lstm_2 (LSTM)                        | (None, 128)     | 80,384  |
| repeat_vector_4 (RepeatVector)       | (None, 28, 128) | 0       |
| lstm_3 (LSTM)                        | (None, 28, 128) | 131,584 |
| time_distributed_4 (TimeDistributed) | (None, 28, 28)  | 3,612   |

Total params: 215,580 (842.11 KB)

Trainable params: 215,580 (842.11 KB)

Non-trainable params: 0 (0.00 B)

## ASSIGNMENT 2

```
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.metrics import confusion_matrix, classification_report
import tensorflow as tf
from tensorflow.keras.datasets import cifar10
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Conv2D, MaxPooling2D, UpSampling2D,
SimpleRNN, LSTM, TimeDistributed, RepeatVector, Flatten
from tensorflow.keras.optimizers import Adam

# Set memory growth for GPU (if available)
gpus = tf.config.experimental.list_physical_devices('GPU')
if gpus:
    try:
        for gpu in gpus:
            tf.config.experimental.set_memory_growth(gpu, True)
    except RuntimeError as e:
        print(e)

# Load and preprocess CIFAR-10 dataset
(x_train, y_train), (x_test, y_test) = cifar10.load_data()
x_train, x_test = x_train.astype('float32') / 255.0, x_test.astype('float32')
/ 255.0

# Reshape data for CNN and RNN/LSTM
x_train_cnn = x_train.reshape(-1, 32, 32, 3)
x_test_cnn = x_test.reshape(-1, 32, 32, 3)
x_train_rnn = x_train.reshape(-1, 32, 32 * 3)
x_test_rnn = x_test.reshape(-1, 32, 32 * 3)

class_names = ['airplane', 'automobile', 'bird', 'cat', 'deer',
               'dog', 'frog', 'horse', 'ship', 'truck']

# CNN autoencoder model
cnn_autoencoder = Sequential([
    Conv2D(32, kernel_size=(3, 3), activation='relu', padding='same',
input_shape=(32, 32, 3)),
    MaxPooling2D(pool_size=(2, 2), padding='same'),
    Conv2D(32, kernel_size=(3, 3), activation='relu', padding='same'),
    MaxPooling2D(pool_size=(2, 2), padding='same'),
    Conv2D(32, kernel_size=(3, 3), activation='relu', padding='same'),
    UpSampling2D(size=(2, 2)),
    Conv2D(32, kernel_size=(3, 3), activation='relu', padding='same'),
    UpSampling2D(size=(2, 2)),
    Conv2D(3, kernel_size=(3, 3), activation='sigmoid', padding='same')
])
```



## ASSIGNMENT 2

```
cnn_autoencoder.compile(optimizer='adam', loss='binary_crossentropy',
metrics=['accuracy'])
cnn_history = cnn_autoencoder.fit(x_train_cnn, x_train_cnn, epochs=10,
batch_size=128, validation_split=0.2)

# LSTM autoencoder model
lstm_autoencoder = Sequential([
    LSTM(128, activation='relu', input_shape=(32, 32 * 3),
return_sequences=False),
    RepeatVector(32),
    LSTM(128, activation='relu', return_sequences=True),
    TimeDistributed(Dense(32 * 3, activation='sigmoid'))
])

lstm_autoencoder.compile(optimizer='adam', loss='binary_crossentropy',
metrics=['accuracy'])
lstm_history = lstm_autoencoder.fit(x_train_rnn, x_train_rnn, epochs=10,
batch_size=128, validation_split=0.2)

# RNN autoencoder model
rnn_autoencoder = Sequential([
    SimpleRNN(128, activation='relu', input_shape=(32, 32 * 3),
return_sequences=False),
    RepeatVector(32),
    SimpleRNN(128, activation='relu', return_sequences=True),
    TimeDistributed(Dense(32 * 3, activation='sigmoid'))
])

rnn_autoencoder.compile(optimizer='adam', loss='binary_crossentropy',
metrics=['accuracy'])
rnn_history = rnn_autoencoder.fit(x_train_rnn, x_train_rnn, epochs=10,
batch_size=128, validation_split=0.2)

# Reconstruct images using all models
cnn_reconstructed = cnn_autoencoder.predict(x_test_cnn)
lstm_reconstructed = lstm_autoencoder.predict(x_test_rnn).reshape(-1, 32, 32,
3)
rnn_reconstructed = rnn_autoencoder.predict(x_test_rnn).reshape(-1, 32, 32, 3)

# Plot original and reconstructed images
n = 10 # number of images to display
plt.figure(figsize=(20, 9))
for i in range(n):
    # Display original
    ax = plt.subplot(4, n, i + 1)
    plt.imshow(x_test[i])
    plt.title("Original")
    plt.axis('off')
```

## ASSIGNMENT 2

```
# Display CNN reconstructed
ax = plt.subplot(4, n, i + 1 + n)
plt.imshow(cnn_reconstructed[i])
plt.title("CNN Reconstructed")
plt.axis('off')

# Display LSTM reconstructed
ax = plt.subplot(4, n, i + 1 + 2 * n)
plt.imshow(lstm_reconstructed[i])
plt.title("LSTM Reconstructed")
plt.axis('off')

# Display RNN reconstructed
ax = plt.subplot(4, n, i + 1 + 3 * n)
plt.imshow(rnn_reconstructed[i])
plt.title("RNN Reconstructed")
plt.axis('off')

plt.show()

# Function to plot class-wise precision, recall, and F1 score
def plot_classification_report(y_true, y_pred, model_name):
    report = classification_report(y_true, y_pred, target_names=class_names,
output_dict=True)
    metrics = ['precision', 'recall', 'f1-score']

    plt.figure(figsize=(18, 6))

    for idx, metric in enumerate(metrics):
        plt.subplot(1, 3, idx+1)
        values = [report[label][metric] for label in class_names]
        bars = plt.barh(class_names, values, color='skyblue',
edgecolor='black')
        plt.title(f'{metric.capitalize()} per Class', fontsize=16)
        plt.xlim(0, 1)
        plt.grid(axis='x', linestyle='--', alpha=0.7)
        plt.xlabel(metric.capitalize(), fontsize=14)
        plt.xticks(fontsize=12)
        plt.yticks(fontsize=12)
        for bar in bars:
            plt.text(bar.get_width() + 0.02, bar.get_y() + bar.get_height()/2,
f'{bar.get_width():.2f}', va='center', fontsize=12)

    plt.suptitle(f'{model_name} Class-wise Performance Metrics', fontsize=18)
    plt.tight_layout(rect=[0, 0, 1, 0.95])
    plt.show()
```

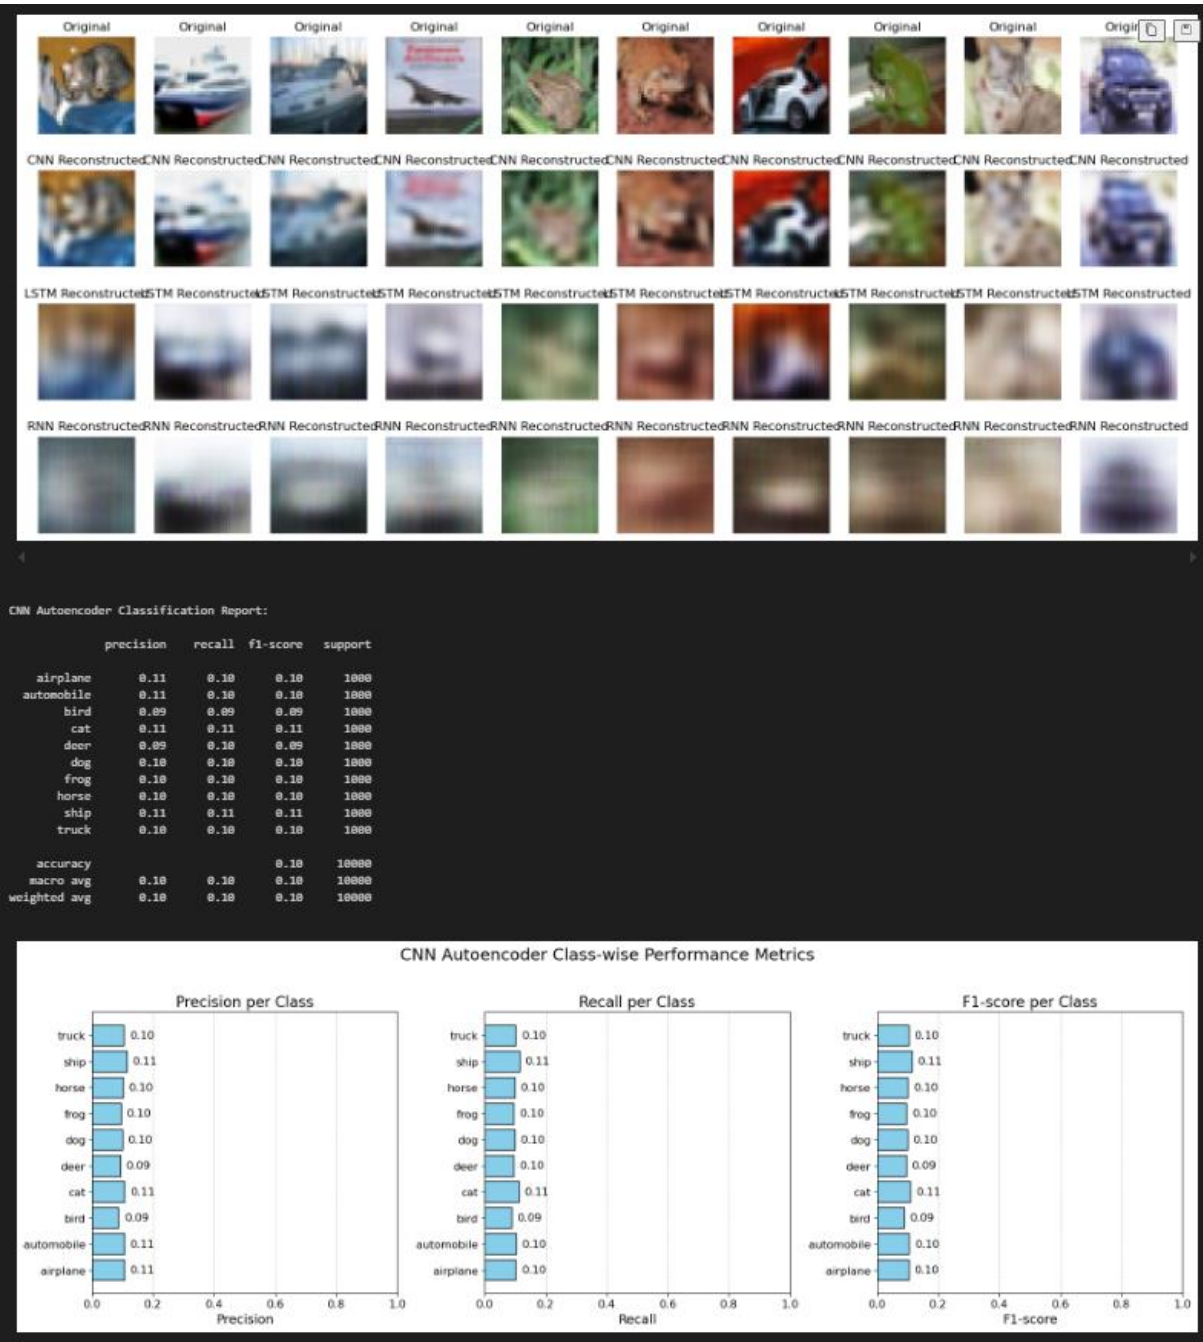
## ASSIGNMENT 2

```
# Simulate predicted labels for evaluation (mock example)
y_pred_cnn = np.random.randint(0, 10, size=len(y_test)) # Replace with actual
predictions
y_pred_lstm = np.random.randint(0, 10, size=len(y_test))
y_pred_rnn = np.random.randint(0, 10, size=len(y_test))

# Plot classification reports for CNN, LSTM, and RNN
def evaluate_model(y_true, y_pred, model_name):
    print(f"\n{model_name} Classification Report:\n")
    print(classification_report(y_true, y_pred, target_names=class_names))
    plot_classification_report(y_true, y_pred, model_name)

evaluate_model(y_test, y_pred_cnn, "CNN Autoencoder")
evaluate_model(y_test, y_pred_lstm, "LSTM Autoencoder")
evaluate_model(y_test, y_pred_rnn, "RNN Autoencoder")
```

ASSIGMENT 2



ASSIGMENT 2

